

Supplementary Materials

We present screenshots of the user interface that was presented to users while performing our three stage online experiment.

Scholar using Natural Language

Scholar is a database with facts about Papers together with their Keyphrases (e.g. Machine Translation), Datasets (e.g. Imagenet), Authors, Conferences and Journals.

Note: Our database is *incomplete*, so results may be fewer than you expect.

E.g: papers by Michael I. Jordan [Run](#)

E.g: keyphrases used by Michael I. Jordan [Run](#)

E.g: How many papers does Michael I. Jordan have ? [Run](#)

Write your natural language query here:
(Try to capitalize noun phrases e.g. *Semantic Parsing*, instead of *semantic parsing*)

Capitalize all names, keywords, years, conferences, paper titles etc.

Execute

Figure 1: Users were presented with example utterances and a text box to enter their own utterance.

Scholar using Natural Language

We recognized the following phrases: *what paper has "michael i. jordan"(AUTHOR) written ?* (All titles, names, years, conferences, keyphrases etc. should be recognized in blue. If not, try capitalizing them)

We have seen a similar query before! paper by "michael i. jordan"(AUTHOR)

Feedback:

- | | |
|---|--|
| <input type="radio"/> Correct | The result answered your question. |
| <input type="radio"/> Wrong Types | The identified names/titles/keyphrase/years (in blue) are not what you intended. |
| <input type="radio"/> Incomplete Result | The result answers your question but is incomplete. For eg. Missing papers or Low number of papers. |
| <input type="radio"/> Wrong Result | The result shows something other than what you wanted, or, the result doesn't make sense. Eg. Authors instead of papers. |
| <input type="radio"/> I can't tell | It looks correct to you, but you're not totally sure. |

Toggle Columns: title abstract numCiting numCitedBy year

Show 10 entries

title

A Variational Principle for Model-based Morphing

Bayesian Nonparametric Inference of Switching Dynamic Linear Models

Convex and Semi-nonnegative Matrix Factorizations: Ding, Li and Jordan
Convex and Semi-nonnegative Matrix Factorizations

Extensions of the Informative Vector Machine

Finite Sample Convergence Rates of Zero-Order Stochastic Optimization Methods

Genome-scale phylogenetic function annotation of large and diverse protein families.

Gradient Descent Only Converges to Minimizers

Learning Graphical Models with Mercer Kernels

Learning in decentralized systems: A nonparametric approach

Learning Spectral Clustering, With Application To Speech Separation

Previous 1 Next

Figure 2: Once the user writes an utterance and pushes execute, they are presented with this screen. First, the identified entities and their types are highlighted to help them decide if the model is receiving the correct inputs. Second, in case the generated (anonymized) query already exists in the training set, it is presented as an alternate utterance, to give users additional confidence about the results that they are seeing. Third, they are presented with 5 feedback options as discussed in the main paper. Finally, they are presented with results and they can toggle columns with additional information.